The background of the slide is a dense field of three-dimensional, light blue numbers (0-9) of various sizes and orientations, creating a sense of depth and data. A solid black rectangular box is positioned on the right side of the slide, containing the title and subtitle text in white.

Singularity and Docker

An introduction to
containerizing a
researcher's workflow

Table of Contents

What are Singularity and Docker?

Pros and Cons to using one over the other

Interactive Example

Where should you go from here?

What is Singularity?



The gold-standard for containerization on HPC clusters



Preferred choice for mission-critical workflows

Used in aerospace, oil and gas, and biotechnology industries



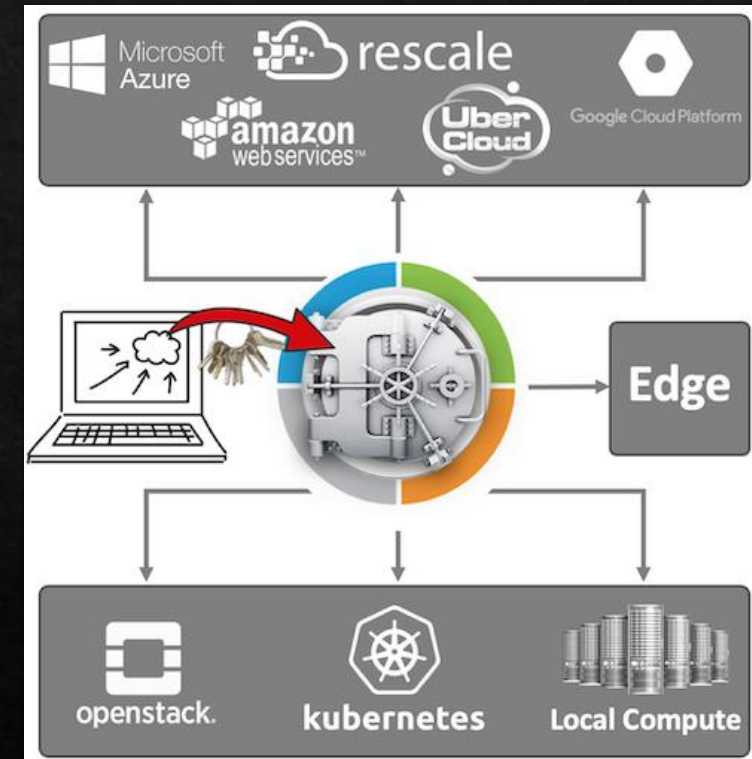
Security first mindset

No one is hacking you!



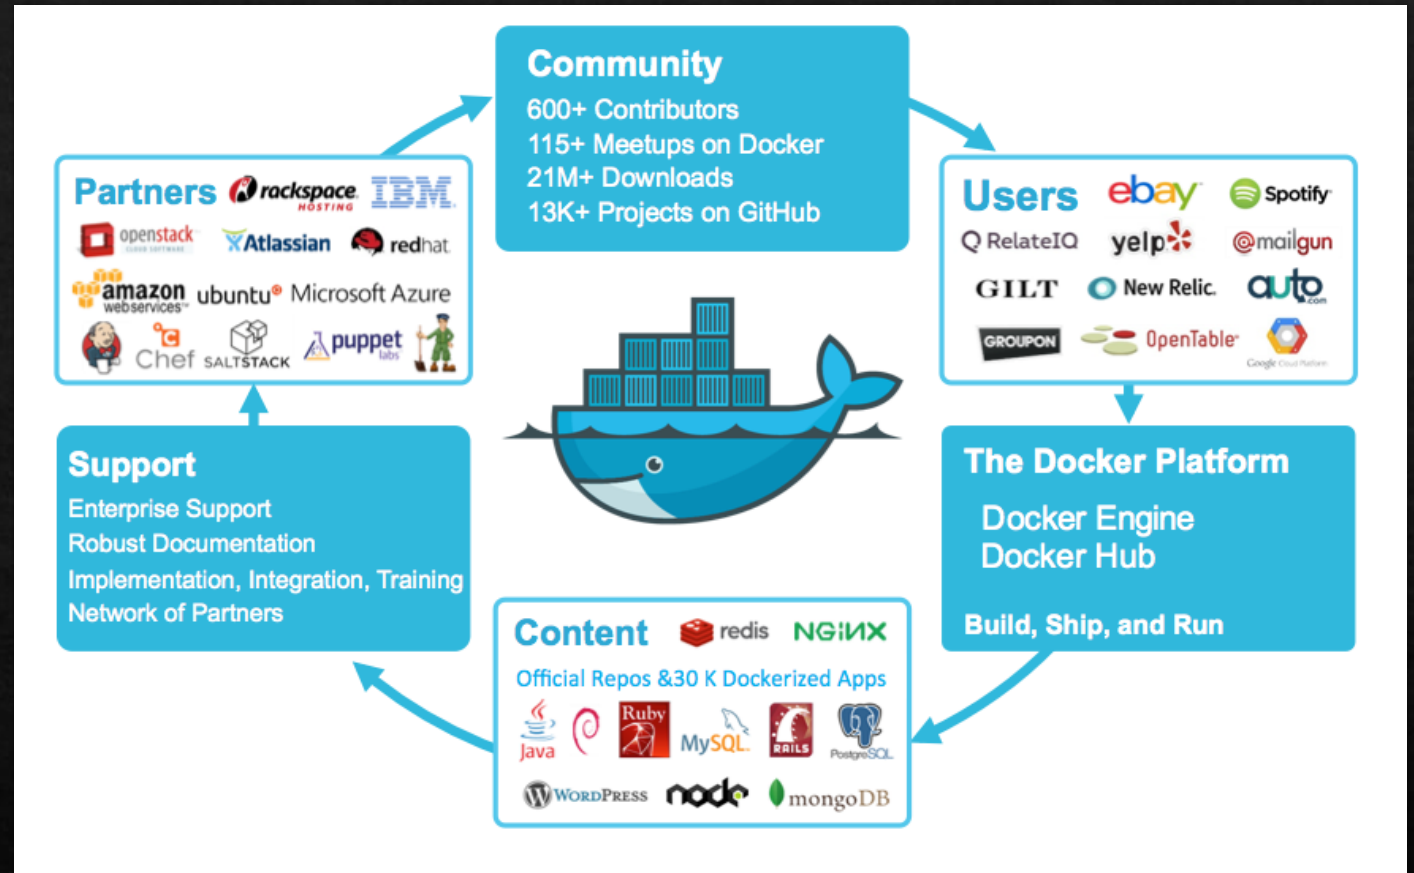
Extremely portable

Can run in any environment that has singularity RE installed



What is Docker?

- ◇ Another popular, ubiquitous containerization software
- ◇ Heavily used in the tech industry
 - ◇ Kubernetes anyone?
- ◇ Extremely adaptable to any kind of workflow
- ◇ Cross-Platform
 - ◇ Works on Linux, MacOS, and Windows



The Pros and Cons of Singularity

The Pros 😊

- ◆ Uber secure
 - ◆ Singularity containers are read-only
- ◆ Space efficient
 - ◆ Containers are compressed to conserve available storage space
- ◆ Installed on Roar
- ◆ Can easily manage system-level dependencies
- ◆ Interoperable with Docker

The Cons ☹️

- ◆ Only supports static environments
- ◆ Only available for use on Linux
 - ◆ Need to use a work-around to build containers on Windows and Mac
- ◆ User community isn't as large as compared to Docker
 - ◆ Finding support can be a little difficult

The Pros and Cons of Docker

The Pros 😊


- ◆ Ubiquitous, virtually used everywhere
- ◆ Cached builds
 - ◆ Start from where you left off if build fails
- ◆ Support dynamic environments
- ◆ Can modify containers after being built from images
- ◆ Multithread container downloads and uploads

The Cons 😞

- ◆ Not installed on the cluster
- ◆ Security is managed by user, not built-in
 - ◆ Container security has become an entire industry because of Docker
- ◆ Not resource conscious
 - ◆ Bloated containers

Interactive Example Time

Set up your build environment!



- ◇ Open the terminal on your computer
- ◇ Create directory titled **VMs**
- ◇ Change into **VMs** directory and create **SINGULARITY** directory
- ◇ In **SINGULARITY** directory, enter the following commands:
 - ◇ `vagrant init sylabs/singularity-3.6-ubuntu-bionic64`
 - ◇ `vagrant up`

Time to build your container!

- ◆ Use the following commands:

- ◆ `vagrant ssh`

- ◆ `git clone https://github.com/NucciTheBoss/iask_onboarding_spring_2021.git`

- ◆ `cd iask_onboarding_spring_2021/day_2`

- ◆ `sudo singularity build image.sif R-4.0.3-rstudio.def`

- ◆ sudo password is “**vagrant**”

- ◆ Now we wait!

A laboratory setting with a graduated cylinder in the foreground, partially filled with a blue liquid. The cylinder has a scale from 0 to 90. In the background, there are several other pieces of glassware, including a round-bottom flask with green liquid, a beaker with yellow liquid, and other flasks, all slightly out of focus. The word "Demonstration!" is overlaid in white text on the right side of the image.

Demonstration!

Where should you go from here?

- ◇ Feel free to play around with my repository:
 - ◇ https://github.com/NucciTheBoss/iask_onboarding_spring_2021
- ◇ Experience is the greatest teacher!
 - ◇ Create your own containers
 - ◇ Play around with containers we already [created](#)
- ◇ Ask for help!
 - ◇ Will, Justin Petucci, and I use Singularity all the time
 - ◇ Add me as a collaborator on a git repository!
 - ◇ I can provide input if you need help!



"That's all Folks!"